



Small Vehicle Image Classification for Pedestrian-Friendly Lanes and Streets




**Deep Learning
Assignment 2**

Highlights



- This project compare 5 difference transfer learning models (VGG16, EfficientNetB7, MobileNetV2, ResNet50, InceptionResNetV2) for 3 classes classification problem (mountain bike, tricycle and tuk tuk)
- You will see what the performance would be for transfer learning
- The interesting question is whether a more complicated model would result a better performance? You can find the answer in this project
- This project applying Grad-CAM to VGG16 model to understand how the model classifies each class

Agenda



01

Introduction

02

**Data Scrapping and
Preparation**

03

**Network
architecture and
Training**

04

Model Evaluation


05

**Ablation Study
(Grad-CAM)**

06

**Discussion and
Conclusion**

Why these small vehicles for image classification?



Cities across the globe become more welcoming to pedestrians and tend to reduce motorized transport amid the self-isolation to curtail the COVID-19 pandemic.

This allows citizens to return to the street at a social distance rather than banning traffic.

Many strategy execution processes for pedestrian-friendly streets are employed nowadays. For example, they provide adequately sized sidewalks and amenities for pedestrians and transit riders. They ensure a good co-existence between motorists and pedestrians.

However, pedestrian's experience should be more pleasant and any inconvenience or danger should be minimized.



Why these small vehicles for image classification?



There are other ways to enhance pedestrian-friendly routes and design to reduce other smaller vehicle-bikes such as mountain bikes, tricycles and Thai tuk-tuk accessing these areas.

Several tech organizations and various disciplines have benefited from CNN and image processing techniques.

Hence, we would like to apply the image classification techniques via CNN to fill this gap and prevent other unwanted vehicles.

ImageNet was also used to explore this because of an extensive image database that has helped advance CNN's research into computer vision and deep learning.

Purpose



The aims of this study are:

- To study multi-label classification task in the effectiveness of small vehicle image classification using five CNN models for pedestrian-friendly lanes and streets *
- To compare the effectiveness of small vehicle image classification using five CNN models for pedestrian-friendly lanes and streets *

* The five CNN models in this study are:

1)VGG-16 2) EfficientNetB7 3) MobileNetV2 4)ResNet50 5)InceptionResNetV2

** These models are randomized to investigate the results according to our class lesson.

Classes



The dataset contains 3 different classes:

- 1. Mountain bike, all terrain-bike, off-roader (class id on IMAGENET: 671)
- 2. Tricycle, trike, velocipede (class id on IMAGENET: 870)
- 3. Tuk tuk (no class on IMAGENET)



Image scraping

```
timeStarted = time.time()
while True:

    imageElement = driver.find_element_by_xpath("//*[id='Sva75c']/div/div/div[3]/div[2]/c-wiz/")
    imageURL = imageElement.get_attribute('src')

    if imageURL != previewImageURL:
        #print("actual URL", imageURL)
        break

    else:
        #making a timeout if the full res image can't be loaded
        currentTime = time.time()

        if currentTime - timeStarted > 10:
            print("Timeout! Will download a lower resolution image and move onto the next one")
            break

#Downloading image
try:
    download_image(imageURL, folder_name, i)
    print("Downloaded element %s out of %s total. URL: %s" % (i, len_containers + 1, imageURL))
except:
    print("Couldn't download an image %s, continuing downloading the next one"%(i))
```

Waiting...

300

Downloaded element 1 out of 301 total. URL: <https://upload.wikimedia.org/wikipedia/commons/6/6e/Ve>
Downloaded element 2 out of 301 total. URL: <https://cdn.britannica.com/69/19469-004-9BCC238A/Veloc>
[es-M-Ives-1869.jpg?w=400&h=300&c=crop](https://cdn.britannica.com/69/19469-004-9BCC238A/Veloces-M-Ives-1869.jpg?w=400&h=300&c=crop)
Downloaded element 3 out of 301 total. URL: <http://img.kansasmemory.org/00617259.jpg>
Downloaded element 4 out of 301 total. URL: <https://merriam-webster.com/assets/mw/images/gallery/ge1698295689-2097-c406a28e99fca87db3c0feac78b61723@1x.jpg>
Downloaded element 5 out of 301 total. URL: <https://upload.wikimedia.org/wikipedia/commons/c/ce/Ve>
Downloaded element 6 out of 301 total. URL: [https://i0.wp.com/ageofrevolution.org/wp-content/uploa](https://i0.wp.com/ageofrevolution.org/wp-content/uploads/2016/07/velo)
[pg?fit=2569%2C1788&ssl=1](https://i0.wp.com/ageofrevolution.org/wp-content/uploads/2016/07/velo)
Downloaded element 7 out of 301 total. URL: https://www.ccpl.org/sites/default/files/Goddard_veloc
Downloaded element 8 out of 301 total. URL: [https://img.pixers.pics/pho_wat\(s3:700/FO/43/64/58/81/](https://img.pixers.pics/pho_wat(s3:700/FO/43/64/58/81/9af53f14244bfd93f7ff.jpg,700,670,cms:2018/10/5bd1b6b8d04b8_220x50-watermark.png,over,480,620,jpg)/)
[9af53f14244bfd93f7ff.jpg,700,670,cms:2018/10/5bd1b6b8d04b8_220x50-watermark.png,over,480,620,jpg\)/](https://img.pixers.pics/pho_wat(s3:700/FO/43/64/58/81/9af53f14244bfd93f7ff.jpg,700,670,cms:2018/10/5bd1b6b8d04b8_220x50-watermark.png,over,480,620,jpg)/)
Downloaded element 9 out of 301 total. URL: <http://dict.drkrok.com/wp-content/uploads/2016/07/velo>
Downloaded element 10 out of 301 total. URL: <https://www.prints-online.com/p/164/latest-style-amer>
[60.jpg.webp](https://www.prints-online.com/p/164/latest-style-amer)
Downloaded element 11 out of 301 total. URL: <https://wikiimg.tojsiabtv.com/wikipedia/commons/thumb>
[de.jpg/225px-The_American_Velocipede.jpg](https://wikiimg.tojsiabtv.com/wikipedia/commons/thumb)



The dataset was mainly scraped from the open dataset and free for use, for instance, pixabay, in which the composition of the images didn't have an identity that can identify the person/people in the images unless the images have the URL reference.

Image preparation

```
for filename in arr:
    file_destination = folder_path + filename
    img = cv2.imread(file_destination)
    try:
        if img.shape[0] < 224 or img.shape[1] < 224 or img.shape[2] != 3:
            print(filename, img.shape)
    except:
        print(filename)
```

```
x = []
y = []

for filename in arr:
    file_destination = folder_path + filename


    img = cv2.imread(file_destination)
    RGB_img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) #CAUTION! SAVED IN RGB NOT BGR
    img = cv2.resize(RGB_img, desired_size)
    #print(img.shape)
    plt.imshow(img)
    x.append(img)
    y.append(1)

print(x)
print(y)
```

```
[array([[124, 111, 77],
        [134, 119, 91],
        [ 92, 76, 53],
        ...,
        [ 81, 71, 59],
        [ 85, 73, 61],
        [ 89, 77, 63]],

       [[ 88, 75, 40],
        [116, 101, 71],
        [ 95, 79, 55],
        ...,
        [ 78, 68, 56],
        [ 83, 71, 59],
        [ 90, 78, 66]],

       [[104, 88, 53],
```



The selected image will be converted to size 224*224 and converted to an array with labels indicating which class they belong to. With this process, the images whose pixel size are smaller than the desired size (224*224) are eliminated. Eventually, there are 200 images/arrays left per class.

Training



In order to show the big picture of training process, we'll introduce the VGG-16 model to be the based model for transfer-learning and fine-tuning with initial set up as shown below.

- Python 3.8.5
- Numpy 1.22.0 fixed random seed = 1234
- Tensorflow 2.7.0 fixed random seed = 5678

Training




1. Load the VGG-16 model as pretrained model with only feature extractor section

```
1 #Load VGG-16 model
2 vgg_extractor = tf.keras.applications.vgg16.VGG16(weights = 'imagenet', include_top = False, input_shape = (224, 224, 3))
3 vgg_extractor.summary()

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 0s 0us/step
58900480/58889256 [=====] - 0s 0us/step
Model: "vgg16"

Layer (type) Output Shape Param #
-----
input_1 (InputLayer) [(None, 224, 224, 3)] 0
block1_conv1 (Conv2D) (None, 224, 224, 64) 1792
block1_conv2 (Conv2D) (None, 224, 224, 64) 36928
block1_pool (MaxPooling2D) (None, 112, 112, 64) 0
block2_conv1 (Conv2D) (None, 112, 112, 128) 73856
block2_conv2 (Conv2D) (None, 112, 112, 128) 147584
block2_pool (MaxPooling2D) (None, 56, 56, 128) 0
block3_conv1 (Conv2D) (None, 56, 56, 256) 295168
block3_conv2 (Conv2D) (None, 56, 56, 256) 590080
block3_conv3 (Conv2D) (None, 56, 56, 256) 590080
block3_pool (MaxPooling2D) (None, 28, 28, 256) 0
block4_conv1 (Conv2D) (None, 28, 28, 512) 1180160
block4_conv2 (Conv2D) (None, 28, 28, 512) 2359808
block4_conv3 (Conv2D) (None, 28, 28, 512) 2359808
block4_pool (MaxPooling2D) (None, 14, 14, 512) 0
block5_conv1 (Conv2D) (None, 14, 14, 512) 2359808
block5_conv2 (Conv2D) (None, 14, 14, 512) 2359808
block5_conv3 (Conv2D) (None, 14, 14, 512) 2359808
block5_pool (MaxPooling2D) (None, 7, 7, 512) 0

Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
```



Training




2. Input the data which were prepared to be 3 classes (0 = mountain bike, 1 = tricycle, 2 = tuk tuk) and split data into train and test sets with test size = 20% as shown below.


```
1 #Train Test Split
2 from sklearn.model_selection import train_test_split
3
4 test_size = 0.2
5 x_train, x_test = train_test_split(x, test_size = test_size, random_state = 3)
6 y_train, y_test = train_test_split(y, test_size = test_size, random_state = 3)
7
8 print(x_train.shape)
9 print(y_train.shape)
10 print(x_test.shape)
11 print(y_test.shape)
```

```
(480, 224, 224, 3)
(480,)
(120, 224, 224, 3)
(120,)
```

```
1 #visualize the first 5 image
2 plt.figure(figsize = (15,5))
3 for i in range(5):
4     plt.subplot(150 + 1 + i).set_title(f'class {y_train[i]}')
5     plt.imshow(x_train[i])
6 plt.show()
```



Training



3. Freeze the layers in model by setting the Trainable to False

```
1 #Recursively freeze all layers in the model first
2 vgg_extractor.trainable = False
3
4 for i, layer in enumerate(vgg_extractor.layers):
5     print(f'Layer {i}: Name = {layer.name}, Trainable = {layer.trainable}')
```



```
Layer 0: Name = input_1, Trainable = False
Layer 1: Name = block1_conv1, Trainable = False
Layer 2: Name = block1_conv2, Trainable = False
Layer 3: Name = block1_pool, Trainable = False
Layer 4: Name = block2_conv1, Trainable = False
Layer 5: Name = block2_conv2, Trainable = False
Layer 6: Name = block2_pool, Trainable = False
Layer 7: Name = block3_conv1, Trainable = False
Layer 8: Name = block3_conv2, Trainable = False
Layer 9: Name = block3_conv3, Trainable = False
Layer 10: Name = block3_pool, Trainable = False
Layer 11: Name = block4_conv1, Trainable = False
Layer 12: Name = block4_conv2, Trainable = False
Layer 13: Name = block4_conv3, Trainable = False
Layer 14: Name = block4_pool, Trainable = False
Layer 15: Name = block5_conv1, Trainable = False
Layer 16: Name = block5_conv2, Trainable = False
Layer 17: Name = block5_conv3, Trainable = False
Layer 18: Name = block5_pool, Trainable = False
```

Training




4. Add new classification head for tuk tuk classification

```
1 x = vgg_extractor.output
2
3 #Add our custom layer(s) to the end of the existing model
4 x = tf.keras.layers.Flatten()(x)
5 x = tf.keras.layers.Dense(512, activation = 'relu')(x)
6 x = tf.keras.layers.Dropout(0.5)(x)
7 new_outputs = tf.keras.layers.Dense(10, activation = 'softmax')(x)
8
9 #construct the main model
10 model = tf.keras.models.Model(inputs = vgg_extractor.inputs, outputs = new_outputs)
11
12 model.summary()
```


```
Model: "model"
Layer (type)                   Output Shape          Param #
-----
input_1 (InputLayer)          [(None, 224, 224, 3)] 0
block1_conv1 (Conv2D)         (None, 224, 224, 64)  1792
block1_conv2 (Conv2D)         (None, 224, 224, 64)  36928
block1_pool (MaxPooling2D)    (None, 112, 112, 64)  0
block2_conv1 (Conv2D)         (None, 112, 112, 128) 73856
block2_conv2 (Conv2D)         (None, 112, 112, 128) 147584
block2_pool (MaxPooling2D)    (None, 56, 56, 128)   0
block3_conv1 (Conv2D)         (None, 56, 56, 256)  295168
block3_conv2 (Conv2D)         (None, 56, 56, 256)  590880
block3_conv3 (Conv2D)         (None, 56, 56, 256)  590880
block3_pool (MaxPooling2D)    (None, 28, 28, 256)   0
block4_conv1 (Conv2D)         (None, 28, 28, 512)  1180160
block4_conv2 (Conv2D)         (None, 28, 28, 512)  2359808
block4_conv3 (Conv2D)         (None, 28, 28, 512)  2359808
block4_pool (MaxPooling2D)    (None, 14, 14, 512)   0
block5_conv1 (Conv2D)         (None, 14, 14, 512)  2359808
block5_conv2 (Conv2D)         (None, 14, 14, 512)  2359808
block5_conv3 (Conv2D)         (None, 14, 14, 512)  2359808
block5_pool (MaxPooling2D)    (None, 7, 7, 512)     0
flatten (Flatten)             (None, 25088)          0
dense (Dense)                  (None, 512)            12845568
dropout (Dropout)             (None, 512)            0
dense_1 (Dense)                (None, 10)             5130

Total params: 27,565,386
Trainable params: 12,850,698
Non-trainable params: 14,714,688
```





Training

Vgg extractor



Our new model



The classification section
for tuk tuk image classifier





Training



5. Model training by transfer-learning

- Loss : Cross Entropy (Sparse Categorical)
- Optimizer : Adam
- Metrics : Accuracy
- Batch Size : 128
- Epoch Number : 20


```
1 model.compile(loss = 'sparse_categorical_crossentropy', optimizer = 'adam', metrics = ['acc'])
2 history = model.fit(x_train_vgg, y_train, batch_size = 128, epochs = 20, verbose = 1, validation_split = 0.2)

Epoch 1/20
3/3 [=====] - 34s 5s/step - loss: 22.7171 - acc: 0.5755 - val_loss: 6.9981 - val_acc: 0.9479
Epoch 2/20
3/3 [=====] - 4s 2s/step - loss: 4.5083 - acc: 0.9349 - val_loss: 8.9988 - val_acc: 0.9583
Epoch 3/20
3/3 [=====] - 4s 2s/step - loss: 3.4484 - acc: 0.9740 - val_loss: 8.9134 - val_acc: 0.9583
Epoch 4/20
3/3 [=====] - 4s 2s/step - loss: 1.7867 - acc: 0.9844 - val_loss: 9.4579 - val_acc: 0.9583
Epoch 5/20
3/3 [=====] - 4s 2s/step - loss: 0.4076 - acc: 0.9870 - val_loss: 9.3418 - val_acc: 0.9688
Epoch 6/20
3/3 [=====] - 4s 2s/step - loss: 0.3216 - acc: 0.9948 - val_loss: 9.4479 - val_acc: 0.9688
Epoch 7/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 9.8066 - val_acc: 0.9792
Epoch 8/20
3/3 [=====] - 4s 2s/step - loss: 0.2044 - acc: 0.9948 - val_loss: 10.2192 - val_acc: 0.9792
Epoch 9/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 11.2178 - val_acc: 0.9688
Epoch 10/20
3/3 [=====] - 4s 2s/step - loss: 0.2509 - acc: 0.9974 - val_loss: 11.7979 - val_acc: 0.9688
Epoch 11/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 12.0260 - val_acc: 0.9688
Epoch 12/20
3/3 [=====] - 4s 2s/step - loss: 0.4837 - acc: 0.9974 - val_loss: 11.9456 - val_acc: 0.9688
Epoch 13/20
3/3 [=====] - 4s 2s/step - loss: 0.2291 - acc: 0.9948 - val_loss: 11.0191 - val_acc: 0.9688
Epoch 14/20
3/3 [=====] - 4s 2s/step - loss: 1.5522e-09 - acc: 1.0000 - val_loss: 10.3491 - val_acc: 0.9792
Epoch 15/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 10.0057 - val_acc: 0.9792
Epoch 16/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 9.7385 - val_acc: 0.9792
Epoch 17/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 9.5249 - val_acc: 0.9792
Epoch 18/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 9.3796 - val_acc: 0.9688
Epoch 19/20
3/3 [=====] - 4s 2s/step - loss: 0.3560 - acc: 0.9974 - val_loss: 9.3767 - val_acc: 0.9688
Epoch 20/20
3/3 [=====] - 4s 2s/step - loss: 0.0000e+00 - acc: 1.0000 - val_loss: 9.3413 - val_acc: 0.9688
```


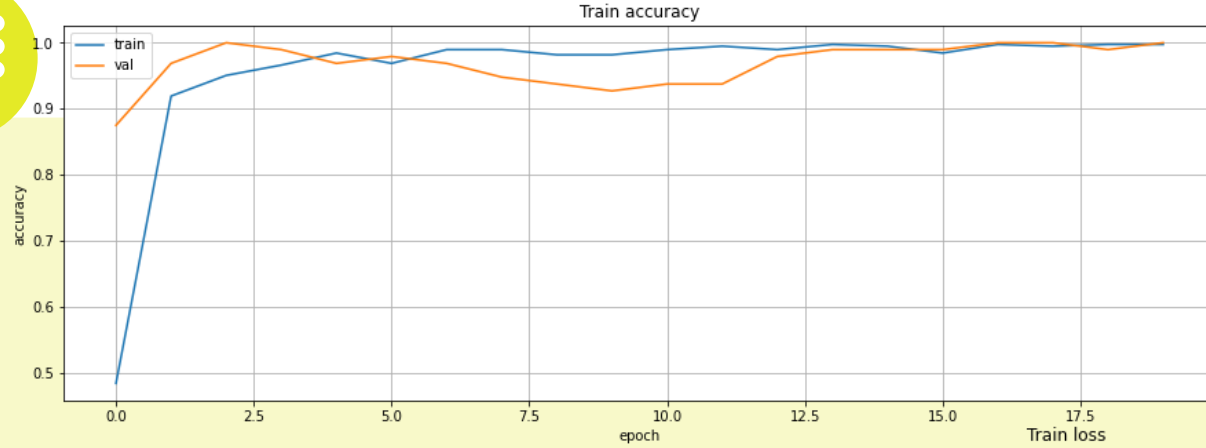
VGG-16 Model




VGG-16 Model





EfficientNetB7 Model




EfficientNetB7 Model





MobileNetV2 Model




MobileNetV2 Model





ResNet50 Model




ResNet50 Model



InceptionResNetV2 Model



InceptionResNetV2 Model



Model Evaluation on Test Set



| Models | Accuracy | Loss |
|-------------------|----------|-------|
| VGG-16 | 0.992 | 2.265 |
| EfficientNetB7 | 0.967 | 2.894 |
| MobileNetV2 | 0.975 | 0.436 |
| ResNet50 | 0.967 | 2.737 |
| InceptionResNetV2 | 0.958 | 0.952 |

Conclusion




According to the experiment, Google was a source for our image scraping and collecting. As the result of five CNN model studies, there were accuracy between 0.958-0.992 and loss between 0.436-2.737. However, the VGG-16 model was shown the best accuracy (0.992) at a loss of 2.265. Thus, we suggest applying the VGG-16 model to classify smaller vehicle bikes for pedestrian-friendly streets

Gradient-weighted Class Activation Mapping (Grad-CAM)



- *Grad-CAM uses the gradients of target flowing into the final convolutional layer to produce a map highlighting regions in the image for predicting the concept* (Ref: [Understand your Algorithm with Grad-CAM | by Daniel Reiff | Towards Data Science](#))
- Applied Grad-CAM to our modified VGG16 model to study which image area the model focus on to classify the mountain bike, tricycle, and tuk-tuk
- The higher heatmap value (colored red in 'jet' colormap) shows the important area for prediction probability



Grad-CAM applied on dog and cat image

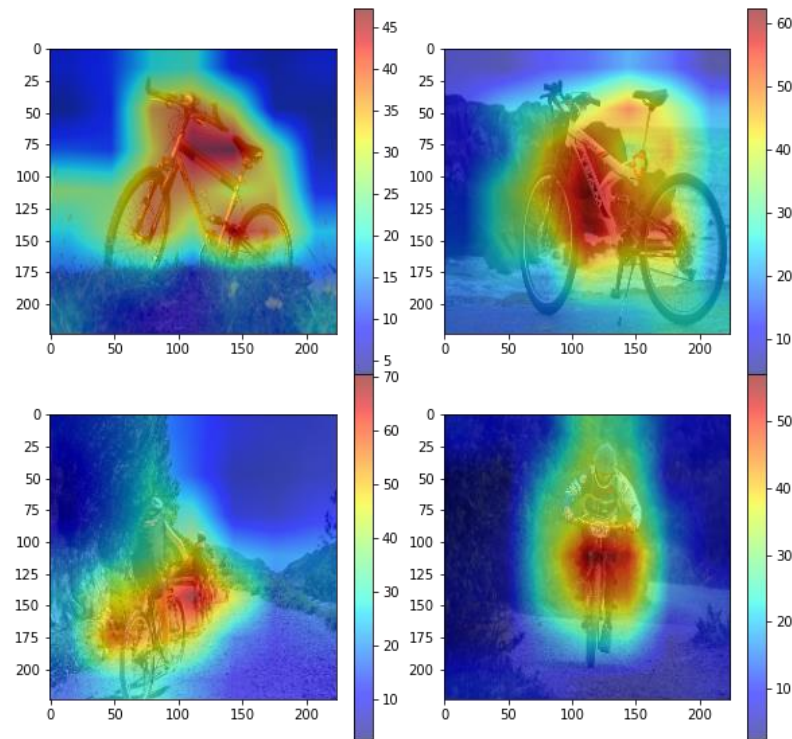
(Ref: [Understand your Algorithm with Grad-CAM | by Daniel Reiff | Towards Data Science](#))

Grad-CAM for Mountain Bike




Finding for Mountain Bike

- The model seems to identify the Mountain Bike by looking at the center core part of the mountain bike




Grad-CAM applied on mountain bike image

Grad-CAM for Tricycle



Finding for **Tricycle**

- The model seems to identify the Tricycle by looking at
 - ☐ the wheel
 - ☐ the handle




Grad-CAM applied on tricycle image

Grad-CAM for Tuk Tuk



Finding for **Tuk Tuk**

- The model seems to identify the Tuk-Tuk by looking at overall unique characteristics of tuk-tuk such as
 - ☐ the front lights
 - ☐ the side
 - ☐ the front windshield



Grad-CAM applied on tuk tuk image

Gradient-weighted Class Activation Mapping (Grad-CAM)



Grad-CAM Conclusion

- Because the differences between the each class are quite obvious
- Applying Grad-CAM to each class shows that the model classifies each class by looking at the class's outstanding features

Grad-CAM Code



```
def alter_model_for_GradCAM(model, last_conv_layer_name):
    last_conv_output = model.get_layer(last_conv_layer_name).output

    old_weights = [x.numpy() for x in model.layers[-1].weights]
    new_config = model.layers[-1].get_config()
    new_config['activation'] = tf.keras.activations.linear
    new_config['name'] = 'prediction_linear'
    out_linear = tf.keras.layers.Dense(**new_config)(model.layers[-2].output)
    out_softmax = tf.keras.activations.softmax(out_linear)

    new_model = tf.keras.Model(inputs=model.inputs, outputs=[out_softmax, out_linear, last_conv_output])
    new_model.layers[-2].set_weights(old_weights)

    return new_model

def my_CNN_GradCAM(model, in_img, class_index):
    in_img = tf.cast(in_img, tf.float32)

    with tf.GradientTape() as tape:
        tape.watch(in_img)
        y_softmax, y_linear, last_conv_activation = model(in_img)
        one_class_score = y_linear[..., class_index]
        gradient = tape.gradient(one_class_score, last_conv_activation)

    gradient = gradient.numpy().squeeze(axis=0)
    alpha = np.mean(gradient, axis=(0,1))

    last_conv_activation = last_conv_activation.numpy().squeeze(axis=0)
    heatmap = np.dot(last_conv_activation, alpha)

    heatmap = np.maximum(0, heatmap)

    return heatmap

new_model = alter_model_for_GradCAM(model, 'block5_pool')
new_model.summary()
```

```
# SHOW 4X4 GRAD-CAM IMAGES FROM SELECTED CLASS

# Input #####
selected_class = 2 #class index {0:'mtbike', 1:'tricycle', 2:'tuktuk'}
#####

import random
import cv2

predict_encode = {0:'mtbike', 1:'tricycle', 2:'tuktuk'}
rand_idx = []

if selected_class == 0:
    xxx = x0
    yyy = y0
elif selected_class == 1:
    xxx = x1
    yyy = y1
elif selected_class == 2:
    xxx = x2
    yyy = y2
else:
    print('Wrong Class Provided')

for i in range(4): #4 images to show
    rand_idx.append(random.randint(0,len(xxx)))

fig = plt.figure(figsize=(10,10))

for i in range(len(rand_idx)):
    print(rand_idx[i])
    x_input = xxx[rand_idx[i]][np.newaxis, ...]
    y_pred = np.argmax(model.predict(x_input))
    print(f'Predicted Class: {predict_encode[y_pred]}')
    print(f'Actual Class: {predict_encode[yyy[rand_idx[i]]]}')

    heatmap = my_CNN_GradCAM(new_model, x_input, y_pred)
    img = xxx[rand_idx[i]]

    ax = fig.add_subplot(2,2,i+1)
    ax.imshow(img)
    alpha = 0.6
    im = ax.imshow(cv2.resize(heatmap, img.shape[:2]), cmap='jet', alpha=alpha)
    plt.colorbar(im, ax=ax)

plt.subplots_adjust(wspace=0.1, hspace=0)
plt.show()
```

Discussion



- The results are quite good with 99.2% accuracy for the VGG-16 model, and all 5 models we selected have more than 95% accuracy. For the next step, we may consider running all models with the same CPU and GPU condition to compare the processing time for each model which will be another key criterion in case of a real implementation.
- Another improvement is to add more classes for small vehicle image classification such as pedestrian, trolley, food truck, etc. To be more beneficial and can be used for more purposes.

References



Python Version: 3.8.5

Python Library

- Matplotlib 3.5.1
- Numpy 1.22.0
- OpenCV 4.5.5
- Tensorflow 2.7.0

Source Code

- Transfer Learning: BADS7604 by Asst. Prof. Thitirat Sriborbornratanakul (CNN2_ex3)
- Grad-CAM: BADS7604 by Asst. Prof. Thitirat Sriborbornratanakul (CNN3_ex2)

Datasets

- Please look at the Citing page

Citing



- For those who want to use the dataset images,
 - ☐ For **mountain bike** dataset, please read and use the image under the **pixabay.com** policy
 - ☐ For **tricycle** dataset, referenced URL are provided in **image_references_tricycle.txt**
 - ☐ For **tuk-tuk** dataset, cite them as a format provided in **image_references_tuktuk.txt**
- For those who want to use any other image but not the dataset images, please reference the image in **bibtex** format

Team Members

6310412018

Pongsarat
Chootai

(25%) Develop image scrapping coding for team, Image Scrapping (Tricycle), Develop original coding for model transfer-learning for team and evaluate VGG-model, Result Discussion for improvement

6310412021

Theethut
Narksenee

(25%)
Image Scrapping (Mountain bike), Image preparation, Transfer Learning of EfficientNetB7 model and InceptionResNetV2 model

6310412024

Saranchai
Angkawinijwong

(25%)
Highlights, Image Scrapping (Tuk Tuk), VGG Transfer Learning, Grad-CAM, Miscellaneous(References, Citing, End Credit)

6420400001

Krittapat
Chuenphitthayavut

(25%)
Introduction composition, Image scaping, Experimental variation on resnet50 and MobileNetV2

End Credit



This project is a part of **Deep Learning** course

Course code: **BADS7604**

Syllabus: Business Analytics and Data Science (**BADS**)

Institute: National Institute of Development Administration (**NIDA**)

Semester: **2/2021**

Thank you for your attention